

(a) (10 points) One case of cache insertion at an index (a set) with exactly one way already occupied. Don't forget to explain the change of LRU bits of both ways in that set.

MEM
EV

Access address: 0000000000000000010001001010011000

Index	V	D	LRU	Tag	Block 0
0	0	0	1		
	0	0	1		
1	0	0	1		
	0	0	1		
2	0	0	1		
	0	0	1		
3	1	0	0	0x000000000000044a	0x0003000000150047
	0	0	1		
4	0	0	1		
	0	0	1		
5	0	0	1		
	0	0	1		
6	0	0	1		
	0	0	1		
7	0	0	1		
	0	0	1		

31 65 32 0

Access address 0111111111111111111111111111011100

Index V D LRU Tag Block 0

0	0	0	1		
	0	0	1		
1	0	0	1		
	0	0	1		
2	0	0	1		
	0	0	1		
3	1	0	1	0x000000000000044a	0x0003000000150047
	1	1	0	0x0000000000001fffff	0x0000000000000000
4	0	0	1		
	0	0	1		
5	0	0	1		
	0	0	1		
6	0	0	1		
	0	0	1		
7	0	0	1		
	0	0	1		

Image1.1.3:accessed memory address & cache state after the instruction

Explanation:

With write allocate, we fetch the block on a **write miss**.

After the write miss at **index 3**, the valid bit(V) and dirty bit(D) are set to 1, and the LRU bit will become 0 to represent that the block is the most recently used, just like what Image1.1.3 shows.

What's more, by Image1.1.2 and 1.1.3, we can see that in the 4th set(with index3), there's **exactly one way already occupied**.

After the **store instruction**, the other way is occupied.

- (b) (10 points) One case of a hit that makes a block become dirty. (Not dirty before that hit)

10220:	02113423	sd x1 40 x2	MEM
--------	----------	-------------	-----

Image1.2.1: the store instruction I find the case

Access address 31 65 32 0
011111111111111111111111111111110011000

Index	V	D	LRU	Tag	Block 0
0	1	1	0	0x0000000001fffffe	0x000000007fffffc0
1	1	1	1	0x0000000001ffffff	0x0000000070000002
1	1	1	1	0x0000000001fffffd	0x0000000070000000
1	1	0	0	0x0000000001fffffe	0x00000000000010270
2	1	1	0	0x0000000001fffffe	0x0000000000000006
2	1	1	1	0x0000000001ffffff	0x0000000090000008
3	1	1	1	0x0000000001fffffd	0x000000007fffff90
3	1	0	0	0x0000000001fffffe	0x000000007fffffc0
4	1	1	1	0x0000000001ffffff	0x0000000000000000
4	1	1	0	0x0000000001fffffd	0x0000000000000006
5	1	1	1	0x0000000001fffffd	0x000000007fffffc0
5	1	1	0	0x0000000001fffffe	0x0000000030000000
6	1	1	0	0x0000000001fffffc	0x000000007fffffd8
6	1	1	1	0x0000000001fffffd	0x0000000060000000
7	1	1	1	0x0000000001fffffc	0x000000007fffffcc
7	1	1	0	0x0000000001fffffd	0x0000000020000006

Image1.2.2: cache state before the instruction

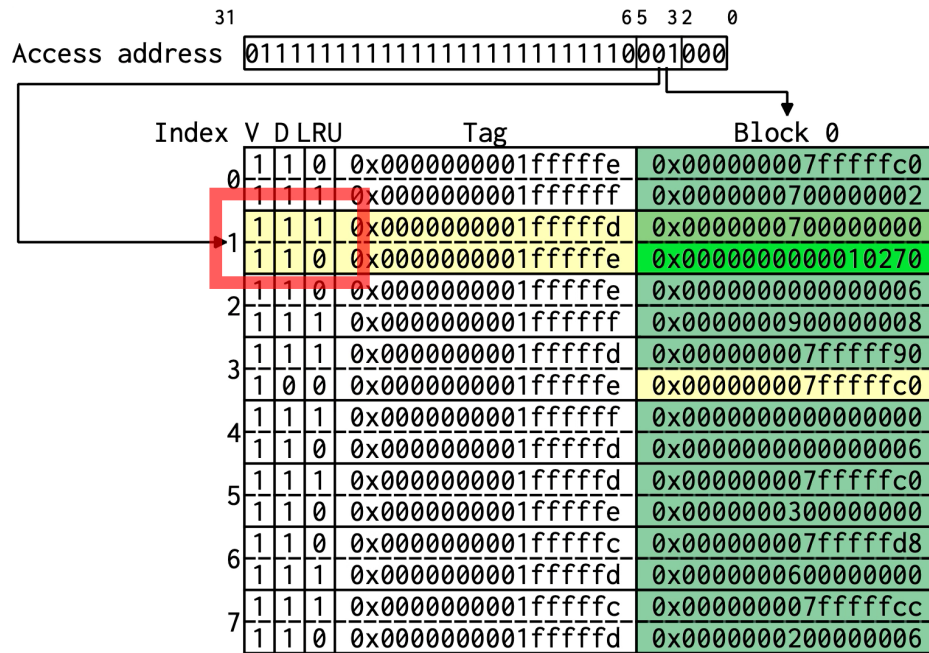


Image1.2.3:accessed memory address & cache state after the instruction

Explanation:

With write allocate, we fetch the block on a **write hit**.

After the write hit at 2nd set(with index 1), the dirty bit(D) is set to 1, just like what Image1.2.2 and 1.2.3 show.

As a result, we can see that the write hit makes a block becomes dirty.

(c) (10 points) One case of a replacement with write-back needed.




Image1.3.1: the store instruction I find the case

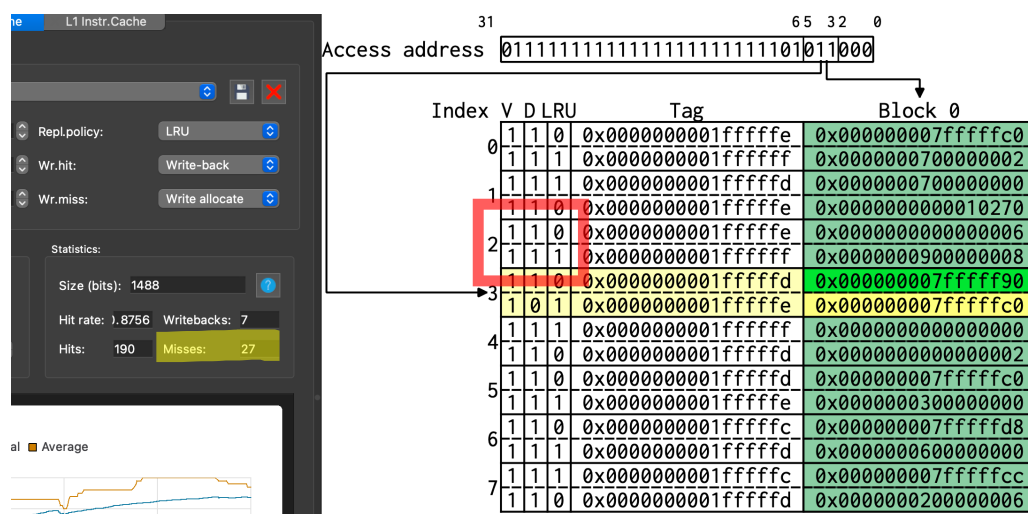


Image1.3.2: cache state before the instruction

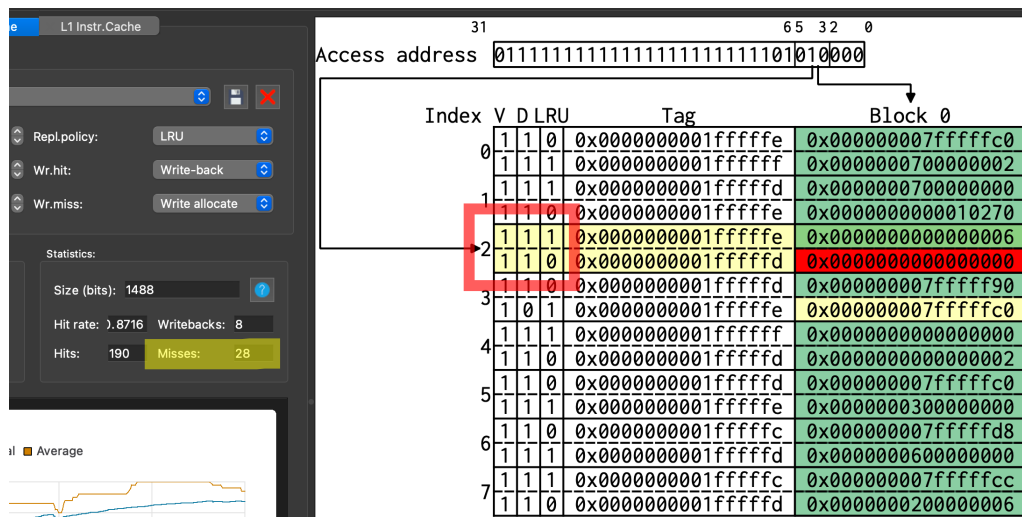


Image1.3.3:accessed memory address & cache state after the instruction

Explanation:

By Image1.3.2 and 1.3.3, we can see that the 28th miss happened after the store instruction, and the LRU bit becomes 0 (LRU bit = 1, originally). What's more, since the dirty bit for the original data is 1, we know that we need to write back the old data to the memory and replace the block with new data.

2. (10 points) Improve the hit rate by designing your own cache. You may adjust the associativity or the cache size. Please give some screenshots to show (i) the cache you design and (ii) the improvement of the hit rate. A brief discussion about why the cache you design makes the hit rate higher is also needed. Note that if you can't get a better hit rate somehow, a reasonable discussion is still needed. Otherwise, you can use the given ELF example that a higher hit rate is definitely possible.

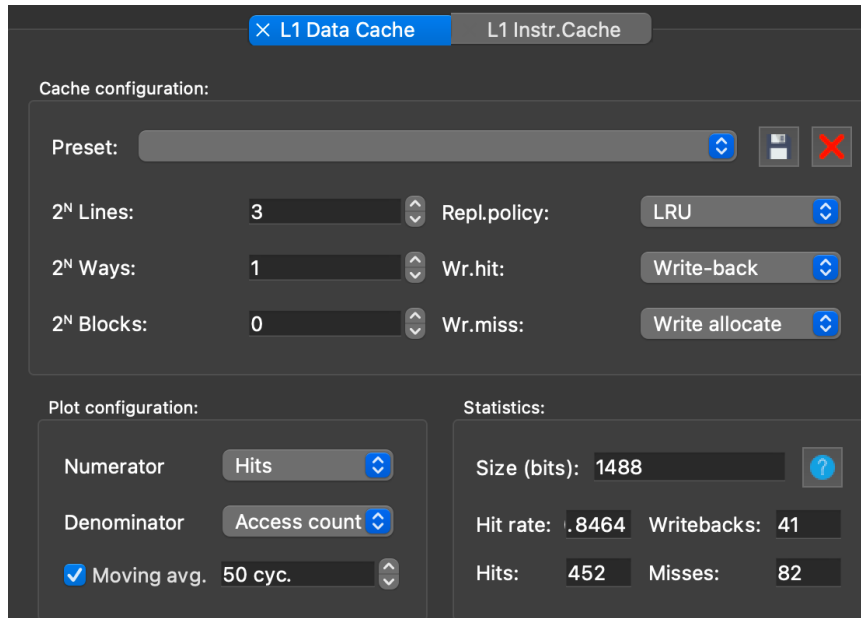


Image2.1: the original cache →84.6%

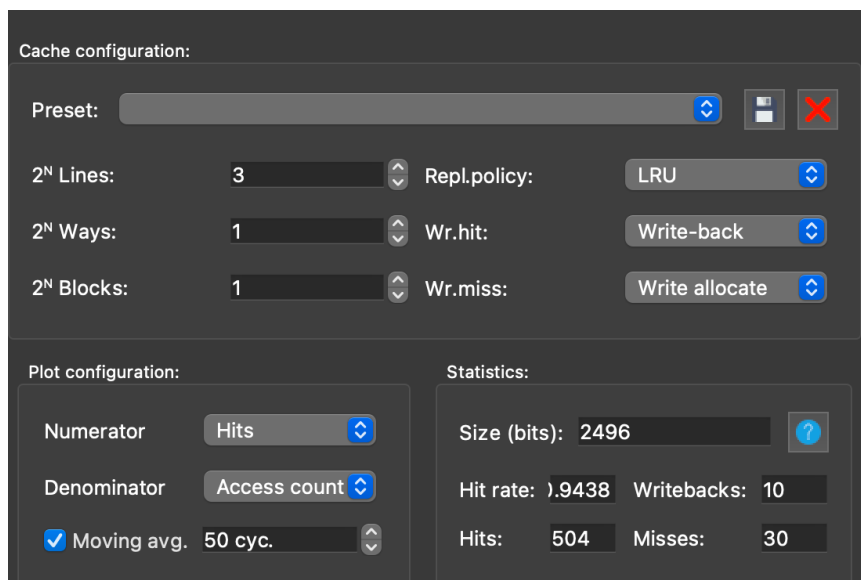


Image2.2: 3-1-1 →94.3%

Cache configuration:

Preset:

2^N Lines: 3 Repl.policy: LRU

2^N Ways: 2 Wr.hit: Write-back

2^N Blocks: 0 Wr.miss: Write allocate

Plot configuration:

Numerator: Hits

Denominator: Access count

☒ Moving avg. 50 cyc.

Statistics:

Size (bits): 3008

Hit rate: .9457 Writebacks: 0

Hits: 505 Misses: 29

Image2.3: 3-2-0→94.6%

Cache configuration:

Preset:

2^N Lines: 4 Repl.policy: LRU

2^N Ways: 1 Wr.hit: Write-back

2^N Blocks: 0 Wr.miss: Write allocate

Plot configuration:

Numerator: Hits

Denominator: Access count

☒ Moving avg. 50 cyc.

Statistics:

Size (bits): 2944

Hit rate: .9288 Writebacks: 9

Hits: 496 Misses: 38

Image2.4: 4-1-0 →92.9%

By image 2.1~2.4, we can see that increase the number of ways and blocks are better than increase the one of lines.

As a result, I will **focus on increase the former two aspect instead of increasing the one of lines to enhance the Hit rate to near 100% later(image 2.5 and image 2.6)**

Cache configuration:

Preset:

2^N Lines: 3 Repl.policy: LRU

2^N Ways: 2 Wr.hit: Write-back

2^N Blocks: 3 Wr.miss: Write allocate

Plot configuration:

Numerator: Hits

Denominator: Access count

☒ Moving avg. 50 cyc.

Statistics:

Size (bits): 17248

Hit rate: 0.9906 Writebacks: 0

Hits: 529 Misses: 5

Image2.5: 3-2-3 →99%

× L1 Data Cache L1 Instr.Cache

Cache configuration:

Preset:

2^N Lines: 3 Repl.policy: LRU

2^N Ways: 3 Wr.hit: Write-back

2^N Blocks: 2 Wr.miss: Write allocate

Plot configuration:

Numerator: Hits

Denominator: Access count

☒ Moving avg. 50 cyc.

Statistics:

Size (bits): 18240

Hit rate: 0.9831 Writebacks: 0

Hits: 525 Misses: 9

Image2.6: 3-3-2 →98%

By image2.5 and 2.6, we can see that the performance of increasing the number of blocks is better than increase that of ways.

Therefore, the cache show by image 2.5 is my final cache.

In my opinion, my coding style contribute to this situation since I use nearly all of the registers in my code and I use stack to store the data before entering a loop or a function almost every time; as a result, increasing the amount of the ways or the blocks is better!

And since the registers I used for store instruction when storing data into the stack are the same type(e.g. temporary register or save register), which means their addresses have little difference, it's better to increase the number of blocks than the ways.

3. (20 points) Suppose a cache receives the following memory addresses for accessing 32-bit single words: (Note: Memory addresses are to the bytes.)

0x74, 0x3c, 0xa4, 0x70, 0x34, 0x08, 0xc4, 0xb8, 0x94, 0xb8

For each reference, identify (i) the index, (ii) the tag, and (iii) whether it is a hit or a miss, for each of the following caches. Assume each cache is initially empty.

- (a) (10 points) A direct-mapped cache with 1-word blocks and a total size of 8 words.

The memory address is 8bit

8/1 = 8 block \rightarrow 3 bit index,

1word/block=4 byte block \rightarrow 2 bit offset

8-3-2=3 \rightarrow 3 bit tag

Word addr	Binary addr	Tag	Index	Hit/Miss
0x74	011 101 00	011	101	Miss
0x3c	001 111 00	001	111	Miss
0xa4	101 001 00	101	001	Miss
0x70	011 100 00	011	100	Miss
0x34	001 101 00	001	101	Miss
0x08	000 010 00	000	010	Miss
0xc4	110 001 00	110	001	Miss
0xb8	101 110 00	101	110	Miss
0x94	100 101 00	100	101	Miss
0xb8	101 110 00	101	110	Hit

- (b) (10 points) A 2-way set associative cache with 2-word blocks and a total size of 16 words.

16/2 = 8 block, 8 block / 2 way = 4 sets \rightarrow 2bit index

2word/block = 8 byte block \rightarrow 3 bit offset

8-2-3=3 \rightarrow 3 bit tag

Word addr	Binary addr	Tag	Index	Hit/Miss
0x74	011 10 100	011	10	Miss
0x3c	001 11 100	001	11	Miss
0xa4	101 00 100	101	00	Miss
0x70	011 10 000	011	10	Hit
0x34	001 10 100	001	10	Miss
0x08	000 01 000	000	01	Miss

0xc4	110 00 100	110	00	Miss
0xb8	101 11 000	101	11	Miss
0x94	100 10 100	100	10	Miss
0xb8	101 11 000	101	11	Hit

4. (10 points) Consider the Hamming single error correcting (SEC) code.

(a) (4 points) Let $A = 10011100$ be an 8-bit data. Show how to get a 12-bit Hamming SEC code word B for A .

8bit $\rightarrow 2^n \geq n + 8 + 1$, $n = 4 \rightarrow 4$ parity bits

__ 1 _ 0 0 1 _ 1 1 0 0

__ 1 _ 0 0 1 _ 1 1 0 0 $\rightarrow p1 = 1$

__ 1 _ 0 0 1 _ 1 1 0 0 $\rightarrow p2 = 1$

__ 1 _ 0 0 1 _ 1 1 0 0 $\rightarrow p4 = 1$

__ 1 _ 0 0 1 _ 1 1 0 0 $\rightarrow p8 = 0$

$B = 111100101100$

(b) (4 points) Suppose C is obtained by inverting the fifth bit from the left of B and hence has a single bit error. Show how to find and correct the single bit error in C .

$B = 111100101100 \rightarrow C = 111110101100$

1111 1010 1100 $\rightarrow H = ___ 1$

1111 1010 1100 $\rightarrow H = __ 0 1$

1111 1010 1100 $\rightarrow H = _ 1 0 1$

1111 1010 1100 $\rightarrow H = 0 1 0 1$

0101 $\rightarrow 5(\text{decimal}) \rightarrow \text{the fifth bit}$

(c) (2 points) What is the minimum number of parity bits required to protect a 256-bit data using the Hamming SEC code?

$2^n \geq n + 256 + 1 = n + 2^8 + 1 \rightarrow n = 9 \rightarrow 9$ parity bits

5. (10 points) The following table gives some information for each of the instruction and data caches attached to each of the two processors, P1 and P2.

	I-cache miss rate	D-cache miss rate	Clock period
P1	5%	8%	0.5 ns
P2	4%	6%	1 ns

Assume each processor has a base CPI of 1 without any memory stalls. Each access to the instruction or data memory takes 50 ns. And 40% of all instructions have data accesses. Note that a base CPI of 1 means that an instruction completes in one clock cycle, unless either the instruction access or the data access causes a cache miss.

(a) (8 points) What is the total CPI for P1 and P2?

$$\text{AMAT} = \text{Hit Time} + \text{Miss Rate} * \text{Miss Penalty}$$

P1 :

$$\text{Hit Time} = 1 \text{ (clock/instruction)} * 0.5 \text{ (ns/clock)} = 0.5\text{ns/instruction}$$

$$\begin{aligned} \text{AMAT} &= 0.5\text{ns} + 50\text{ns} * 0.05 + 50\text{ns} * 0.08 * 0.4 \\ &= 0.5 + 2.5 + 1.6 = 4.6\text{ns} \end{aligned}$$

Note: I-cache miss may happen in every instructions!!Remember not to multiply 0.4

$$\text{CPI} = (4.6\text{ns /instruction}) / (0.5\text{ns/clock}) = \underline{9.2\text{clock/instruction}}$$

P2:

$$\text{Hit Time} = 1 \text{ (clock/instruction)} * 1 \text{ (ns/clock)} = 1\text{ns/instruction}$$

$$\begin{aligned} \text{AMAT} &= 1\text{ns} + 50\text{ns} * 0.04 + 50\text{ns} * 0.06 * 0.4 \\ &= 1 + 2 + 1.2 = \text{ns} \end{aligned}$$

$$\text{CPI} = (4.2\text{ns /instruction}) / (1\text{ns/clock}) = \underline{4.2\text{clock/instruction}}$$

(b) (2 points) Which processor is faster if they have the same amount of instructions?

Time cost per instruction:

$$\text{P1: } 9.2 * 0.5 \text{ ns} = 4.6\text{ns}$$

$$\text{P2: } 4.2 * 1\text{ns} = 4.2\text{ns}$$

$$4.2 < 4.6, 4.6 - 4.2 = 0.4$$

P2 is faster than P1 by 0.4ns

6. (20 points) The following data constitute a stream of virtual byte addresses generated by a processor.

0x4a90, 0x210c, 0xa3b5, 0x0e18, 0x904f

With **4 KB** pages, a **four-entry fully associative TLB**, and **approximate LRU replacement**, assume that the main memory has adequate space to accommodate the requested physical page brought from the disk without page replacement. Therefore, you can choose any page number as long as it is not used in the page table.

The initial TLB and page table states are as below.

TLB			
Valid	Tag	Physical Page Number	Reference (Used)
1	0x5	11	1
1	0x0	5	0
1	0xa	3	0
0	0x3	6	1

Page table		
Index	Valid	Physical Page Number or in Disk
0	1	5
1	0	Disk
2	0	Disk
3	1	6
4	1	9
5	1	11
6	0	Disk
7	1	4
8	0	Disk
9	0	Disk
a	1	3
b	1	12

For each access shown above, list whether (1) the access is a hit or miss in the TLB, (2) whether the access causes a page fault, and (3) the updated state of the TLB.

Page number(4bit)	Offset(12bit)
-------------------	---------------

address	Hit/Miss Page Fault	TLB				
		Valid	Tag	PPN	reference	Replacement pointer
0x4a90 (0100 1010 1001 0000)	Miss in TLB Hit in PT	1	0x5	11	1	v
		1	0x0	5	0	
		1	0xa	3	0	
		1	0x4	9	1	

address	Hit/Miss Page Fault	TLB				
		Valid	Tag	PPN	reference	Replacement pointer
0x210c (0010 0001 0000 1100)	Miss in TLB and PT Page fault	1	0x5	11	0	
		1	0x2	1 (random Choose)	1	
		1	0xa	3	0	v
		1	0x4	9	1	

address	Hit/Miss Page Fault	TLB				
		Valid	Tag	PPN	reference	Replacement pointer
0xa3b5 (1010 0011 1011 0101)	Hit in TLB	1	0x5	11	0	
		1	0x2	1	1	
		1	0xa	3	1	v
		1	0x4	9	1	

address	Hit/Miss Page Fault	TLB				
		Valid	Tag	PPN	reference	Replacement pointer
0x0e18 (0000 1110 0001 1000)	Miss in TLB Hit in PT	1	0x0	5	1	
		1	0x2	1	1	v
		1	0xa	3	0	
		1	0x4	9	0	

address	Hit/Miss Page Fault	TLB				
		Valid	Tag	PPN	reference	Replacement pointer
0x904f (1001 0000 0100 1111)	Miss in TLB and PT Page fault	1	0x0	5	1	
		1	0x2	1	0	
		1	0x9	2 (random Choose)	1	
		1	0x4	9	0	v